

practice-assignment-4

May 4, 2025

```
[1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
[2]: df=pd.read_csv("BostonHousing.csv")
df
```

```
[2]:
```

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	\
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	
..
501	0.06263	0.0	11.93	0	0.573	6.593	69.1	2.4786	1	273	
502	0.04527	0.0	11.93	0	0.573	6.120	76.7	2.2875	1	273	
503	0.06076	0.0	11.93	0	0.573	6.976	91.0	2.1675	1	273	
504	0.10959	0.0	11.93	0	0.573	6.794	89.3	2.3889	1	273	
505	0.04741	0.0	11.93	0	0.573	6.030	80.8	2.5050	1	273	
	ptratio	b	lstat	medv							
0	15.3	396.90	4.98	24.0							
1	17.8	396.90	9.14	21.6							
2	17.8	392.83	4.03	34.7							
3	18.7	394.63	2.94	33.4							
4	18.7	396.90	5.33	36.2							
..							
501	21.0	391.99	9.67	22.4							
502	21.0	396.90	9.08	20.6							
503	21.0	396.90	5.64	23.9							
504	21.0	393.45	6.48	22.0							
505	21.0	396.90	7.88	11.9							

[506 rows x 14 columns]

```
[3]: print(df.head())
      print(df.tail())
      print(df.info())
      print(df.describe())
```

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	\
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	

	b	lstat	medv
0	396.90	4.98	24.0
1	396.90	9.14	21.6
2	392.83	4.03	34.7
3	394.63	2.94	33.4
4	396.90	5.33	36.2

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	\
501	0.06263	0.0	11.93	0	0.573	6.593	69.1	2.4786	1	273	21.0	
502	0.04527	0.0	11.93	0	0.573	6.120	76.7	2.2875	1	273	21.0	
503	0.06076	0.0	11.93	0	0.573	6.976	91.0	2.1675	1	273	21.0	
504	0.10959	0.0	11.93	0	0.573	6.794	89.3	2.3889	1	273	21.0	
505	0.04741	0.0	11.93	0	0.573	6.030	80.8	2.5050	1	273	21.0	

	b	lstat	medv
501	391.99	9.67	22.4
502	396.90	9.08	20.6
503	396.90	5.64	23.9
504	393.45	6.48	22.0
505	396.90	7.88	11.9

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 506 entries, 0 to 505

Data columns (total 14 columns):

#	Column	Non-Null Count	Dtype
0	crim	506 non-null	float64
1	zn	506 non-null	float64
2	indus	506 non-null	float64
3	chas	506 non-null	int64
4	nox	506 non-null	float64
5	rm	501 non-null	float64
6	age	506 non-null	float64
7	dis	506 non-null	float64
8	rad	506 non-null	int64
9	tax	506 non-null	int64
10	ptratio	506 non-null	float64

```

11  b          506 non-null    float64
12  lstat      506 non-null    float64
13  medv       506 non-null    float64

```

dtypes: float64(11), int64(3)

memory usage: 55.5 KB

None

	crim	zn	indus	chas	nox	rm \
count	506.000000	506.000000	506.000000	506.000000	506.000000	501.000000
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284341
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.705587
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.884000
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208000
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.625000
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000

	age	dis	rad	tax	ptratio	b \
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	68.574901	3.795043	9.549407	408.237154	18.455534	356.674032
std	28.148861	2.105710	8.707259	168.537116	2.164946	91.294864
min	2.900000	1.129600	1.000000	187.000000	12.600000	0.320000
25%	45.025000	2.100175	4.000000	279.000000	17.400000	375.377500
50%	77.500000	3.207450	5.000000	330.000000	19.050000	391.440000
75%	94.075000	5.188425	24.000000	666.000000	20.200000	396.225000
max	100.000000	12.126500	24.000000	711.000000	22.000000	396.900000

	lstat	medv
count	506.000000	506.000000
mean	12.653063	22.532806
std	7.141062	9.197104
min	1.730000	5.000000
25%	6.950000	17.025000
50%	11.360000	21.200000
75%	16.955000	25.000000
max	37.970000	50.000000

```

[4]: print(df.shape)
      print(df.size)

```

```

(506, 14)
7084

```

```

[5]: print(df.isna())
      print(df.isna().sum())

```

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax \
0	False	False	False	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False	False	False	False

```

2    False False False False False False False False False False False
3    False False False False False False False False False False False
4    False False False False False False False False False False False
..    ...    ...    ...    ...    ...    ...    ...    ...    ...    ...
501  False False False False False False False False False False False
502  False False False False False False False False False False False
503  False False False False False False False False False False False
504  False False False False False False False False False False False
505  False False False False False False False False False False False

```

```

      ptratio      b  lstat  medv
0      False False False False
1      False False False False
2      False False False False
3      False False False False
4      False False False False
..      ...    ...    ...    ...
501  False False False False
502  False False False False
503  False False False False
504  False False False False
505  False False False False

```

[506 rows x 14 columns]


```

crim      0
zn        0
indus     0
chas      0
nox       0
rm        5
age       0
dis       0
rad       0
tax       0
ptratio   0
b         0
lstat     0
medv      0
dtype: int64

```

```
[6]: df.boxplot()
```

```
[6]: <Axes: >
```



```
[7]: Q1 = df['medv'].quantile(0.25)
Q3 = df['medv'].quantile(0.75)
IQR = Q3 - Q1
Lower_limit = Q1 - 1.5 * IQR
Upper_limit = Q3 + 1.5 * IQR
print(f'Q1 = {Q1}, Q3 = {Q3}, IQR = {IQR}, Lower_limit = {Lower_limit}, Upper_limit = {Upper_limit}')
```

```
Q1 = 17.025, Q3 = 25.0, IQR = 7.975000000000001, Lower_limit = 5.0624999999999996, Upper_limit = 36.962500000000006
```

```
[8]: outliers_medv=[]
for i in df.medv:
    if i<Lower_limit or i>Upper_limit:
        outliers_medv.append(i)
print("outliers are",outliers_medv)
```

```
outliers are [38.7, 43.8, 41.3, 50.0, 50.0, 50.0, 50.0, 37.2, 39.8, 37.9, 50.0, 37.0, 50.0, 42.3, 48.5, 50.0, 44.8, 50.0, 37.6, 46.7, 41.7, 48.3, 42.8, 44.0, 50.0, 43.1, 48.8, 50.0, 43.5, 45.4, 46.0, 50.0, 37.3, 50.0, 50.0, 50.0, 50.0, 50.0, 5.0, 5.0]
```

```
[9]: df[df.medv<Lower_limit].index
```

```
[9]: Index([398, 405], dtype='int64')
```

```
[10]: outlier_indices=df[(df.medv<Lower_limit)|(df.medv>Upper_limit)].index
df1=df.drop(outlier_indices)
df1
```

```
[10]:
```

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	\
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	
..	
501	0.06263	0.0	11.93	0	0.573	6.593	69.1	2.4786	1	273	
502	0.04527	0.0	11.93	0	0.573	6.120	76.7	2.2875	1	273	
503	0.06076	0.0	11.93	0	0.573	6.976	91.0	2.1675	1	273	
504	0.10959	0.0	11.93	0	0.573	6.794	89.3	2.3889	1	273	
505	0.04741	0.0	11.93	0	0.573	6.030	80.8	2.5050	1	273	

	ptratio	b	lstat	medv
0	15.3	396.90	4.98	24.0
1	17.8	396.90	9.14	21.6
2	17.8	392.83	4.03	34.7
3	18.7	394.63	2.94	33.4
4	18.7	396.90	5.33	36.2
..
501	21.0	391.99	9.67	22.4
502	21.0	396.90	9.08	20.6
503	21.0	396.90	5.64	23.9
504	21.0	393.45	6.48	22.0
505	21.0	396.90	7.88	11.9

[466 rows x 14 columns]

```
[11]: df1.boxplot()
```

```
[11]: <Axes: >
```

```
[12]: outliers_medv=[]
for i in df1.medv:
    if i<Lower_limit or i>Upper_limit:
        outliers_medv.append(i)
print("outliers are",outliers_medv)
```

outliers are []

```
[13]: df1
```

```
[13]:      crim    zn  indus  chas    nox    rm  age    dis  rad  tax  \
0    0.00632  18.0   2.31    0  0.538  6.575  65.2  4.0900    1  296
1    0.02731   0.0   7.07    0  0.469  6.421  78.9  4.9671    2  242
2    0.02729   0.0   7.07    0  0.469  7.185  61.1  4.9671    2  242
3    0.03237   0.0   2.18    0  0.458  6.998  45.8  6.0622    3  222
4    0.06905   0.0   2.18    0  0.458  7.147  54.2  6.0622    3  222
..      ...    ...    ...    ...    ...    ...    ...    ...    ...
501  0.06263   0.0  11.93    0  0.573  6.593  69.1  2.4786    1  273
502  0.04527   0.0  11.93    0  0.573  6.120  76.7  2.2875    1  273
503  0.06076   0.0  11.93    0  0.573  6.976  91.0  2.1675    1  273
504  0.10959   0.0  11.93    0  0.573  6.794  89.3  2.3889    1  273
505  0.04741   0.0  11.93    0  0.573  6.030  80.8  2.5050    1  273

      ptratio    b  lstat  medv
0      15.3  396.90   4.98  24.0
1      17.8  396.90   9.14  21.6
2      17.8  392.83   4.03  34.7
3      18.7  394.63   2.94  33.4
4      18.7  396.90   5.33  36.2
..      ...    ...    ...    ...
501     21.0  391.99   9.67  22.4
502     21.0  396.90   9.08  20.6
503     21.0  396.90   5.64  23.9
504     21.0  393.45   6.48  22.0
505     21.0  396.90   7.88  11.9
```

[466 rows x 14 columns]

```
[14]: X = df.drop(['medv'], axis = 1)
      Y = df['medv']
```

```
[15]: print(X)
      print(Y)
```

```
      crim    zn  indus  chas    nox    rm  age    dis  rad  tax  \
0    0.00632  18.0   2.31    0  0.538  6.575  65.2  4.0900    1  296
1    0.02731   0.0   7.07    0  0.469  6.421  78.9  4.9671    2  242
2    0.02729   0.0   7.07    0  0.469  7.185  61.1  4.9671    2  242
3    0.03237   0.0   2.18    0  0.458  6.998  45.8  6.0622    3  222
4    0.06905   0.0   2.18    0  0.458  7.147  54.2  6.0622    3  222
..      ...    ...    ...    ...    ...    ...    ...    ...    ...
501  0.06263   0.0  11.93    0  0.573  6.593  69.1  2.4786    1  273
502  0.04527   0.0  11.93    0  0.573  6.120  76.7  2.2875    1  273
503  0.06076   0.0  11.93    0  0.573  6.976  91.0  2.1675    1  273
504  0.10959   0.0  11.93    0  0.573  6.794  89.3  2.3889    1  273
```

```
505  0.04741    0.0  11.93    0  0.573  6.030  80.8  2.5050    1  273
```

```
      ptratio      b  lstat
0      15.3  396.90  4.98
1      17.8  396.90  9.14
2      17.8  392.83  4.03
3      18.7  394.63  2.94
4      18.7  396.90  5.33
..      ...      ...
501     21.0  391.99  9.67
502     21.0  396.90  9.08
503     21.0  396.90  5.64
504     21.0  393.45  6.48
505     21.0  396.90  7.88
```

```
[506 rows x 13 columns]
```

```
0      24.0
1      21.6
2      34.7
3      33.4
4      36.2
...
501     22.4
502     20.6
503     23.9
504     22.0
505     11.9
```

```
Name: medv, Length: 506, dtype: float64
```

```
[49]: from sklearn.model_selection import train_test_split
      xtrain, xtest, ytrain, ytest = train_test_split(X, Y, test_size=0.
      ↪2, random_state = 0)
```

```
[51]: import sklearn
      from sklearn.linear_model import LinearRegression
      lm = LinearRegression()
```

```
[53]: from sklearn.impute import SimpleImputer
      imputer = SimpleImputer(strategy='mean') # You can choose other strategies like
      ↪ 'median' or 'most_frequent'
      X = imputer.fit_transform(X)

      # 2. Split the data after imputation
      xtrain, xtest, ytrain, ytest = train_test_split(X, Y, test_size=0.2,
      ↪ random_state=0)

      # 3. Create and train the model
```



```
lm = LinearRegression()
model = lm.fit(xtrain, ytrain)
```

```
[55]: ytrain_pred = lm.predict(xtrain)
      ytest_pred = lm.predict(xtest)
      ytrain_pred
```

```
[55]: array([32.54370319, 21.932219 , 27.54552964, 23.62527133,  6.578493 ,
          14.97841931, 22.21772491, 29.1656419 , 33.23633623, 13.13958076,
          20.27455237, 20.68413823, 12.66852204, 23.3697738 ,  5.02810654,
          19.82447761,  9.43610482, 44.62829537, 30.78647869, 12.51813161,
          17.73377577, 21.38820917, 23.6372432 , 20.44934681, 34.99179512,
          13.88126045, 21.08693233, 35.14506971, 19.42370996, 13.14651005,
          14.07551034, 23.11554812, 14.35905012, 31.26696558, 25.30600475,
          15.4186878 , 24.22888048,  9.38631784, 14.92650845, 20.8121123 ,
          32.71894972, 27.98210777, 25.59531311, 15.57980306, 31.11707339,
          27.96643339, 13.98777371,  7.63296162, 28.439442 , 25.35431753,
           4.50921003, 28.38300415, 16.98778052, 29.76898387, 20.46579326,
          15.92063503, 17.90107052, 12.73790055,  8.74450855, 19.2209131 ,
          34.49718351, 32.92991179, 23.69727449, 19.56029786, 22.84821335,
          26.87821751, 21.82625519, 17.07317532, 32.05572536, 10.93501836,
          19.43965502, 32.49620518, 18.84673792, 15.94631864, 18.64040032,
          14.44326484, 24.59149605, 24.32961808, 16.64743526, 13.33257237,
          20.22362526, 25.14078145, 17.16142814, 24.71659709, 20.81755107,
          27.98436411, 35.61445424, 16.64616352, 11.82105901, 34.81722148,
          30.82293852, 20.74723986, 39.53769412, 28.93973655, 29.14758957,
          17.3850748 , 26.81988505, 39.99994502, 28.7436029 , 16.45233953,
          37.42520467, 35.50043267, 13.4301604 , 29.16282174, 21.61445353,
          24.33573863, 21.42879278, 23.70414569, 27.75720913, 29.66634603,
          14.17071791, 26.10123154, 23.301429 , 12.78870854, 13.69996698,
          25.25219339, 19.33992149, 30.53939588, 10.99425038, 23.59381765,
          16.97608003, 16.95715843, 22.63737653, 21.67809796, 11.7771548 ,
          25.20102569, 28.70665016, 20.15560512, 12.58969204, 25.491827 ,
          25.94150764, 25.1012182 , 23.54704847, 26.75021658, 16.60998619,
          21.8143152 , 36.13010572, 20.99151854, 35.84068525, 25.7240647 ,
          21.53807905, 15.86301742, 31.29104399, 21.25252923, 27.76555697,
          14.82988215, 32.22181362, 13.98471864,  1.73549252, 19.34419812,
          14.27907805, 37.50921956, 15.73128883, 14.43177617, 27.30962563,
          23.26035056, 18.51338996, 30.56987939, 27.27579387, 27.29636247,
          24.83120316, 24.16491603, 25.02281109, 11.15470273, 20.75645507,
          13.62885662, 17.19251126, 12.73588136, 28.36489633, 14.93423297,
          16.28759868, 28.71821577, 14.92423512, 21.2574851 , 12.84398221,
          13.88240604, 22.63704415, 21.20611254, 14.77408684, 20.9615425 ,
          16.95022031, 24.57066259, 12.55888031, 34.75093634, 12.01928324,
          43.11627076, 31.24437499, 35.2678383 , 21.46189294, 15.76203608,
          26.55522047, 29.48442291, 14.09084558, 26.53043308, 37.03279546,
          17.66599331, 10.58127564, 34.12553809, 35.60724777, 18.30871709,
```

22.55422403, 17.99645545, 24.37480858, 19.53272297, 27.33759414,
-4.37710883, 20.55848845, 35.20657023, 36.61182186, 25.09838298,
27.19384405, 20.76511499, 20.59850139, 15.88235358, 20.7109413 ,
20.56015402, 27.91074931, 19.64144048, 7.44448041, 16.38421739,
32.4058924 , 35.20068077, 17.52510488, 18.73818756, 23.38476377,
6.9101272 , 21.474421 , 24.03945664, 16.48224996, 18.39477039,
21.89450064, 27.6087426 , 25.48978313, 36.99207872, 15.44312984,
28.60930065, 25.91061939, 22.289867 , 38.68527907, 20.87019495,
23.42254099, 22.8715894 , 12.46815126, 20.32563551, 33.59413957,
24.81623268, 18.00822181, 33.52610644, 21.60855452, 28.34626703,
32.25732348, 36.72521596, 22.24201483, 24.00412219, 22.43880668,
31.79684758, 22.36905758, 18.85033348, 21.83585813, 28.2440051 ,
22.53534761, 21.83903797, 17.02520576, 17.48214385, 16.95645501,
17.43479517, 16.49537109, 31.59311175, 23.80259153, 17.56107622,
19.79933557, 33.6701167 , 13.95877659, 24.9589198 , 17.37744496,
30.49453397, 29.97850997, 22.60575292, 20.826961 , 34.973321 ,
22.63480818, 32.88008718, 20.76686679, 31.40022292, 30.90899101,
37.56345649, 26.84896675, 21.94002219, 28.69666315, 16.17783819,
26.97516809, 21.11061277, 30.45452101, 9.96317971, 30.88843471,
5.84554101, 15.63338238, 18.16309422, 35.40089044, 32.05626117,
11.05941886, 13.29457062, 21.62687263, 34.40247671, 18.64931021,
19.18796195, 14.99721921, 25.77438304, 41.10507292, 25.03604626,
42.00238887, 24.9324625 , 22.30586831, 12.24950075, 12.02438944,
14.16234864, 18.47767294, 3.08442408, 27.4989249 , 26.07818635,
40.99536581, 21.09643555, 21.1699912 , 34.05630845, 33.45094622,
9.72577257, 24.74893414, 43.31925292, 16.92814859, 17.89933566,
25.52211908, 18.42060434, 6.13576017, 19.33435897, 34.89718181,
16.2265251 , 23.02431044, 13.589263 , 24.53484354, 18.77565936,
17.32092633, 18.75893027, 33.11765205, 19.46432946, 30.73066746,
32.75844496, 41.26823949, 19.17037266, 16.62926097, 37.52686446,
17.99662752, 9.42808957, 15.16127034, 24.93324217, 19.68070885,
16.62013838, 27.45888678, 12.96205602, 5.83839312, 19.02447968,
9.86977983, 28.0976158 , 4.52913758, 29.189982 , 32.14578868,
22.15775469, 16.77843434, 18.07456012, 20.72185108, 33.5981505 ,
27.78362104, 19.53364617, 20.72032337, 6.66638479, 28.92171723,
24.6198858 , 22.18005872, 13.66189533, 25.80558744, 19.360516 ,
8.83678532, 26.69250461, 16.1742696 , 31.37220918, 32.61583684,
25.4444746 , 18.52294419, 30.58689102, 21.5617584 , 25.26864563,
25.90791006, 31.60220704, 24.53163973, 34.45150746, 17.12157442,
19.72932588, 18.55636829, 40.97364198, 25.1446076 , 19.51587917,
33.32027502, 23.79132039, 18.47117877, 23.25099877])

```
[57]: from sklearn.metrics import mean_squared_error, r2_score
mse = mean_squared_error(ytrain, ytrain_pred)
print("The model performance for training set")
print("-----")
print('MSE is {}'.format(mse))
```

```

print("\n")
# model evaluation for testing set
#y_test_predict = lin_model.predict(X_test)
mse = mean_squared_error(ytest, ytest_pred)

print("The model performance for testing set")
print("-----")
print('MSE is {}'.format(mse))
print("\n\n\n")
rmse = (np.sqrt(mean_squared_error(ytrain, ytrain_pred)))
r2 = r2_score(ytrain, ytrain_pred)
print("The model performance for training set")
print("-----")
print('RMSE is {}'.format(rmse))
print('R2 score is {}'.format(r2))
print("\n")
# model evaluation for testing set
#y_test_predict = lin_model.predict(X_test)
rmse = (np.sqrt(mean_squared_error(ytest, ytest_pred)))
r2 = r2_score(ytest, ytest_pred)
print("The model performance for testing set")
print("-----")
print('RMSE is {}'.format(rmse))
print('R2 score is {}'.format(r2))

```

```

The model performance for training set
-----
MSE is 19.391606535694894

```

```

The model performance for testing set
-----
MSE is 33.46497598891468

```

```

The model performance for training set
-----
RMSE is 4.4035901870740535
R2 score is 0.7722485406464419

```

```

The model performance for testing set
-----
RMSE is 5.784892046435671
R2 score is 0.5890259426012614

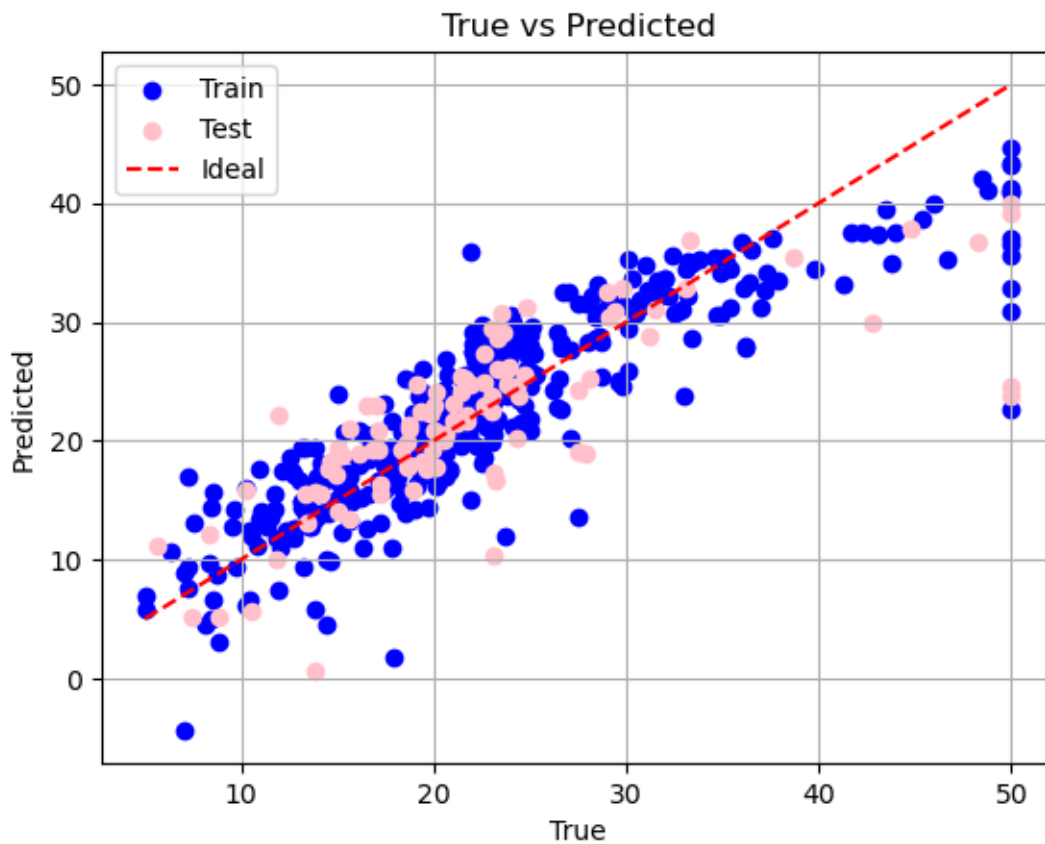
```

```
[59]: import matplotlib.pyplot as plt

plt.scatter(ytrain, ytrain_pred, color='blue', label='Train')
plt.scatter(ytest, ytest_pred, color='pink', label='Test')

# Plot ideal line y = x
min_val = min(min(ytrain), min(ytest))
max_val = max(max(ytrain), max(ytest))
plt.plot([min_val, max_val], [min_val, max_val], 'r--', label='Ideal')

plt.xlabel('True')
plt.ylabel('Predicted')
plt.title('True vs Predicted')
plt.legend()
plt.grid(True)
plt.show()
```



```
[ ]:
```