

practice-assignment-5

May 4, 2025

```
[3]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
[5]: vm=pd.read_csv("Social_Network_Ads.csv")
vm
```

```
[5]:
```

	Age	EstimatedSalary	Purchased
0	19	19000	0
1	35	20000	0
2	26	43000	0
3	27	57000	0
4	19	76000	0
..
395	46	41000	1
396	51	23000	1
397	50	20000	1
398	36	33000	0
399	49	36000	1

[400 rows x 3 columns]

```
[9]: print(vm.head())
print(vm.tail())
print(vm.info())
print(vm.describe())
print(vm.isnull().sum())
```

	Age	EstimatedSalary	Purchased
0	19	19000	0
1	35	20000	0
2	26	43000	0
3	27	57000	0
4	19	76000	0

	Age	EstimatedSalary	Purchased
395	46	41000	1
396	51	23000	1

```

397    50          20000          1
398    36          33000          0
399    49          36000          1
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 3 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Age             400 non-null   int64
1   EstimatedSalary 400 non-null   int64
2   Purchased       400 non-null   int64
dtypes: int64(3)
memory usage: 9.5 KB
None

```

	Age	EstimatedSalary	Purchased
count	400.000000	400.000000	400.000000
mean	37.655000	69742.500000	0.357500
std	10.482877	34096.960282	0.479864
min	18.000000	15000.000000	0.000000
25%	29.750000	43000.000000	0.000000
50%	37.000000	70000.000000	0.000000
75%	46.000000	88000.000000	1.000000
max	60.000000	150000.000000	1.000000
Age	0		
EstimatedSalary	0		
Purchased	0		


```

dtype: int64

```

```
[12]: vm.boxplot()
```

```
[12]: <Axes: >
```



```
[14]: X=vm.drop(['Purchased'],axis=1)
      Y=vm['Purchased']
```

```
[16]: print(X)
      print(Y)
```

	Age	EstimatedSalary
0	19	19000
1	35	20000
2	26	43000
3	27	57000
4	19	76000
..
395	46	41000
396	51	23000
397	50	20000
398	36	33000
399	49	36000

[400 rows x 2 columns]

0	0
1	0

```
2      0
3      0
4      0
```

```
..
395    1
396    1
397    1
398    0
399    1
```

Name: Purchased, Length: 400, dtype: int64

```
[18]: #Split the dataset into training and testing datasets
from sklearn.model_selection import train_test_split
xtrain, xtest, ytrain, ytest = train_test_split(X, Y, test_size=0.
↪2, random_state = 0)
```

```
[26]: from sklearn.linear_model import LogisticRegression
lr=LogisticRegression()
```

```
[30]: lr.fit(xtrain,ytrain)
```

```
[30]: LogisticRegression()
```

```
[32]: y_pred=lr.predict(xtest)
```

```
[34]: print(xtrain)
print("-----\n")
print(xtest)
print("-----\n")
print(ytrain)
print("-----\n")
print(ytest)
print("-----\n")
print(y_pred)
```

	Age	EstimatedSalary
336	58	144000
64	59	83000
55	24	55000
106	26	35000
300	58	38000
..
323	48	30000
192	29	43000
117	36	52000
47	27	54000
172	26	118000

[320 rows x 2 columns]

	Age	EstimatedSalary
132	30	87000
309	38	50000
341	35	75000
196	30	79000
246	35	50000
..
14	18	82000
363	42	79000
304	40	60000
361	53	34000
329	47	107000

[80 rows x 2 columns]

336	1
64	0
55	0
106	0
300	1
..	..
323	1
192	0
117	0
47	0
172	0

Name: Purchased, Length: 320, dtype: int64

132	0
309	0
341	0
196	0
246	0
..	..
14	0
363	0
304	0
361	1
329	1

Name: Purchased, Length: 80, dtype: int64

[0 0 0 0 0 0 0 1 0 1 0 0 0 0 0 0 0 0 1 0 0 1 0 1 0 1 0 0 0 0 0 0 1 0 0 0 0

```

0 0 1 0 0 0 0 1 0 0 1 0 1 1 0 0 0 1 0 0 0 0 0 1 0 0 0 1 0 0 0 0 1 0 0 0
0 0 1 0 1 1]

```

```

[36]: from sklearn.metrics import
      accuracy_score, classification_report, confusion_matrix, precision_score, recall_score, f1_score

```

```

[44]: cm=confusion_matrix(ytest,y_pred)
      cm

```

```

[44]: array([[56,  2],
            [ 5, 17]], dtype=int64)

```

```

[46]: ac=accuracy_score(ytest,y_pred)
      ac

```

```

[46]: 0.9125

```

```

[50]: tn,fp,fn,tp=confusion_matrix(ytest,y_pred).ravel()
      print(tn)
      print(tp)
      print(fn)
      print(fp)

```

```

56
17
5
2

```

```

[52]: err_rate=1-ac
      print(err_rate)

```

```

0.087500000000000002

```

```

[54]: print(classification_report(ytest,y_pred))

```

	precision	recall	f1-score	support
0	0.92	0.97	0.94	58
1	0.89	0.77	0.83	22
accuracy			0.91	80
macro avg	0.91	0.87	0.89	80
weighted avg	0.91	0.91	0.91	80

```

[56]: print(precision_score(ytest,y_pred))
      print(recall_score(ytest,y_pred))
      print(f1_score(ytest,y_pred))

```

0.8947368421052632
0.7727272727272727
0.8292682926829268

[]: